

LOAN STATUS PREDICTION

cris

7/29/2024

STEP 1: READIND AND UNDERSTANDING THE DATA

```
# Clear the memory and check the path for the data
rm(list=ls())
getwd()

## [1] "C:/Users/hp/Desktop/Projects/loan prediction using machine learning"

# Change the path to the desired one
setwd("C:\\Users\\hp\\Desktop\\Projects\\loan prediction using machine
learning")
getwd()

## [1] "C:/Users/hp/Desktop/Projects/loan prediction using machine learning"

# Read the data and print the first 6 variables
data = read.csv("loan-train.csv")
head(data)
```

##	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	
##	1	LP001002	Male	No	0	Graduate	No
##	2	LP001003	Male	Yes	1	Graduate	No
##	3	LP001005	Male	Yes	0	Graduate	Yes
##	4	LP001006	Male	Yes	0	Not Graduate	No
##	5	LP001008	Male	No	0	Graduate	No
##	6	LP001011	Male	Yes	2	Graduate	Yes

##	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	
##	1	0	NA	360	1
##	2	1508	128	360	1
##	3	0	66	360	1
##	4	2358	120	360	1

```

## 5          0          141          360          1
Urban
## 6          4196          267          360          1
Urban
##   Loan_Status
## 1           Y
## 2           N
## 3           Y
## 4           Y
## 5           Y
## 6           Y

# Import the necessary libraries
library(dplyr)

## Warning: package 'dplyr' was built under R version 4.1.3

##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

library(tidyverse)

## Warning: package 'tidyverse' was built under R version 4.1.3
## Warning: package 'ggplot2' was built under R version 4.1.3
## Warning: package 'tibble' was built under R version 4.1.3
## Warning: package 'tidyr' was built under R version 4.1.3
## Warning: package 'readr' was built under R version 4.1.3
## Warning: package 'purrr' was built under R version 4.1.3
## Warning: package 'stringr' was built under R version 4.1.3
## Warning: package 'forcats' was built under R version 4.1.3
## Warning: package 'lubridate' was built under R version 4.1.3

## -- Attaching core tidyverse packages ----- tidyverse
2.0.0 --
## v forcats    1.0.0      v readr      2.1.4
## v ggplot2    3.4.2      v stringr    1.5.0

```

```

## v lubridate 1.9.2      v tibble      3.2.1
## v purrr      1.0.1      v tidyr      1.3.0

## -- Conflicts -----
tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()      masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all
conflicts to become errors

library(ggplot2)
library(stargazer)

## Warning: package 'stargazer' was built under R version 4.1.2

##
## Please cite as:
##
## Hlavac, Marek (2022). stargazer: Well-Formatted Regression and Summary
Statistics Tables.
## R package version 5.2.3. https://CRAN.R-project.org/package=stargazer

# view the structure of our dataset
str(data)

## 'data.frame':    614 obs. of  13 variables:
## $ Loan_ID          : chr  "LP001002" "LP001003" "LP001005" "LP001006" ...
## $ Gender            : chr  "Male" "Male" "Male" "Male" ...
## $ Married           : chr  "No" "Yes" "Yes" "Yes" ...
## $ Dependents        : chr  "0" "1" "0" "0" ...
## $ Education         : chr  "Graduate" "Graduate" "Graduate" "Not Graduate"
...
## $ Self_Employed     : chr  "No" "No" "Yes" "No" ...
## $ ApplicantIncome    : int   5849 4583 3000 2583 6000 5417 2333 3036 4006
12841 ...
## $ CoapplicantIncome: num    0 1508 0 2358 0 ...
## $ LoanAmount         : int   NA 128 66 120 141 267 95 158 168 349 ...
## $ Loan_Amount_Term   : int   360 360 360 360 360 360 360 360 360 ...
## $ Credit_History     : int    1 1 1 1 1 1 1 0 1 1 ...
## $ Property_Area      : chr   "Urban" "Rural" "Urban" "Urban" ...
## $ Loan_Status        : chr    "Y" "N" "Y" "Y" ...

# summarize the dataset in a table format
stargazer(data, type = "text")

##
## =====
## Statistic          N      Mean    St. Dev.   Min      Max
## -----
## ApplicantIncome    614  5,403.459  6,109.042   150    81,000
## CoapplicantIncome  614  1,621.246  2,926.248   0.000 41,667.000
## LoanAmount         592   146.412   85.587     9      700

```

```
## Loan_Amount_Term  600  342.000  65.120  12  480
## Credit_History    564  0.842  0.365  0  1
## -----
```

The first row 'N' shows that there are missing values in the dataset. However, ApplicantIncome and CoapplicantIncome seems to be right skewed due to their large standard deviation. Credit history has only two values, 0 and 1, hence it's not an integer but rather a factor with two levels, i.e., the applicant has either taken a loan before or not. Loan_Amount_Term appears to be symmetrical.

```
summary(data)
```

```
##      Loan_ID          Gender          Married          Dependents
## Length:614      Length:614      Length:614      Length:614
## Class :character Class :character Class :character Class :character
## Mode  :character Mode  :character Mode  :character Mode  :character
##
##
##
##      Education      Self_Employed      ApplicantIncome CoapplicantIncome
## Length:614      Length:614      Min.   : 150      Min.   :  0
## Class :character Class :character 1st Qu.: 2878      1st Qu.:  0
## Mode  :character Mode  :character Median : 3812      Median : 1188
##                                     Mean  : 5403      Mean  : 1621
##                                     3rd Qu.: 5795      3rd Qu.: 2297
##                                     Max.   :81000      Max.   :41667
##
##      LoanAmount      Loan_Amount_Term Credit_History      Property_Area
## Min.   :  9.0      Min.   : 12      Min.   :0.0000      Length:614
## 1st Qu.:100.0      1st Qu.:360      1st Qu.:1.0000      Class :character
## Median :128.0      Median :360      Median :1.0000      Mode  :character
## Mean   :146.4      Mean   :342      Mean   :0.8422
## 3rd Qu.:168.0      3rd Qu.:360      3rd Qu.:1.0000
## Max.   :700.0      Max.   :480      Max.   :1.0000
## NA's   :22      NA's   :14      NA's   :50
##      Loan_Status
## Length:614
## Class :character
## Mode  :character
##
##
##
```

STEP 2: DATA CLEANING. First we check for missing values.

```
# Lets check for missing values
any(is.na(data))
```

```
## [1] TRUE
```

We have missing values. lets check how many missing values do we have

```
sum(is.na(data))  
## [1] 86  
  
# now, we check if we have empty entries and convert them to NA  
colSums(is.na(data)|data == "")  
  
##           Loan_ID           Gender           Married           Dependents  
##           0           13           3           15  
##           Education       Self_Employed       ApplicantIncome       CoapplicantIncome  
##           0           32           0           0  
##           LoanAmount       Loan_Amount_Term       Credit_History       Property_Area  
##           22           14           50           0  
##           Loan_Status  
##           0  
  
# gender, married, dependents,Self_Employed, LoanAmount, Loan_Amount_Term and  
# Credit_History have emty entries. we then convert them to NA  
data[data == ""] = NA  
sum(is.na(data))  
## [1] 149
```

We then handle the missing values

```
data2 = na.omit(data)  
sum(is.na(data2))  
## [1] 0
```

We have no missing valus, nesxt we deal with datatypes

```
# first, we convert character variable to factor form  
data3 = as.data.frame(unclass(data2), stringsAsFactors = TRUE)  
data3$Credit_History = as.factor(data3$Credit_History)  
str(data3)  
  
## 'data.frame':    480 obs. of  13 variables:  
## $ Loan_ID      : Factor w/ 480 levels "LP001003","LP001005",...: 1 2 3  
4 5 6 7 8 9 10 ...  
## $ Gender       : Factor w/ 2 levels "Female","Male": 2 2 2 2 2 2 2 2  
2 2 ...  
## $ Married     : Factor w/ 2 levels "No","Yes": 2 2 2 1 2 2 2 2 2 2  
...  
## $ Dependents  : Factor w/ 4 levels "0","1","2","3+": 2 1 1 1 3 1 4 3  
2 3 ...  
## $ Education   : Factor w/ 2 levels "Graduate","Not Graduate": 1 1 2  
1 1 2 1 1 1 1 ...  
## $ Self_Employed : Factor w/ 2 levels "No","Yes": 1 2 1 1 2 1 1 1 1 1  
...  
...
```

```
## $ ApplicantIncome : int 4583 3000 2583 6000 5417 2333 3036 4006 12841
3200 ...
## $ CoapplicantIncome: num 1508 0 2358 0 4196 ...
## $ LoanAmount : int 128 66 120 141 267 95 158 168 349 70 ...
## $ Loan_Amount_Term : int 360 360 360 360 360 360 360 360 360 360 ...
## $ Credit_History : Factor w/ 2 levels "0","1": 2 2 2 2 2 1 2 2 2 ...
## $ Property_Area : Factor w/ 3 levels "Rural","Semiurban",...: 1 3 3 3 3
3 2 3 2 3 ...
## $ Loan_Status : Factor w/ 2 levels "N","Y": 1 2 2 2 2 2 1 2 1 2 ...
```

Now lets check for duplicates

```
duplicates = duplicated(data3)|duplicated(data3, fromlast = TRUE)
any(duplicates)

## [1] FALSE
```

The data is now clean for analysis


STEP 3: DATA VISUALIZATION (EDA)

```
# Combine columns for income
data4 = data3 %>% mutate(income = data3$ApplicantIncome +
data3$CoapplicantIncome)
head(data4)
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed
ApplicantIncome						
## 1	LP001003	Male	Yes	1	Graduate	No
4583						
## 2	LP001005	Male	Yes	0	Graduate	Yes
3000						
## 3	LP001006	Male	Yes	0	Not Graduate	No
2583						
## 4	LP001008	Male	No	0	Graduate	No
6000						
## 5	LP001011	Male	Yes	2	Graduate	Yes
5417						
## 6	LP001013	Male	Yes	0	Not Graduate	No
2333						
## CoapplicantIncome						
LoanAmount						
Loan_Amount_Term						
Credit_History						
Property_Area						
## 1			1508		128	
					360	
Rural						1
## 2			0		66	
					360	
Urban						1
## 3			2358		120	
					360	
Urban						1
## 4			0		141	
					360	
Urban						1
## 5			4196		267	
					360	
						1


```
Urban
## 6          1516          95          360          1
Urban
##  Loan_Status income
## 1           N   6091
## 2           Y   3000
## 3           Y   4941
## 4           Y   6000
## 5           Y   9613
## 6           Y   3849
```

```
# Now Lets check the distribution for income
hist(data4$income)
```



We can see that income is not normally distributed, ie, it is right skewed. Most of the data tend to lie below 20000

```
# Visualizing the boxplot
boxplot(data4$income)
```



The variable income has also some outliers. lets try to handle the case for outliers and normality by using logarithmic transformation.

```
data5 = data4 %>% mutate(income_transformed = log(income))
head(data5)
```

##	Loan_ID	Gender	Married	Dependents	Education	Self_Employed
## 1	LP001003	Male	Yes	1	Graduate	No
4583						
## 2	LP001005	Male	Yes	0	Graduate	Yes
3000						
## 3	LP001006	Male	Yes	0	Not Graduate	No
2583						
## 4	LP001008	Male	No	0	Graduate	No
6000						
## 5	LP001011	Male	Yes	2	Graduate	Yes
5417						
## 6	LP001013	Male	Yes	0	Not Graduate	No
2333						
##	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Property_Area	
## 1	1508	128	360	1		
Rural						
## 2	0	66	360	1		
Urban						
## 3	2358	120	360	1		



```

Urban
## 4          0          141          360          1
Urban
## 5         4196          267          360          1
Urban
## 6         1516           95          360          1
Urban
##   Loan_Status income income_transformed
## 1           N   6091             8.714568
## 2           Y   3000             8.006368
## 3           Y   4941             8.505323
## 4           Y   6000             8.699515
## 5           Y   9613             9.170872
## 6           Y   3849             8.255569


```

check the previous distribution

```
hist(data5$income_transformed)
```




```
boxplot(data5$income_transformed)
```




Income now looks to be normal with a reduced number of outliers lets check some other sitributions

```
pie(table(data5$Gender))
```




```
pie(table(data5$Married))
```



```
table(data5$Education, data5$Loan_Status)
```

```
##  
##           N    Y  
## Graduate  112 271  
## Not Graduate  36  61
```

```
barplot(table(data5$Education, data5$Loan_Status))
```



```
table(data5$Dependents, data5$Loan_Status)
```

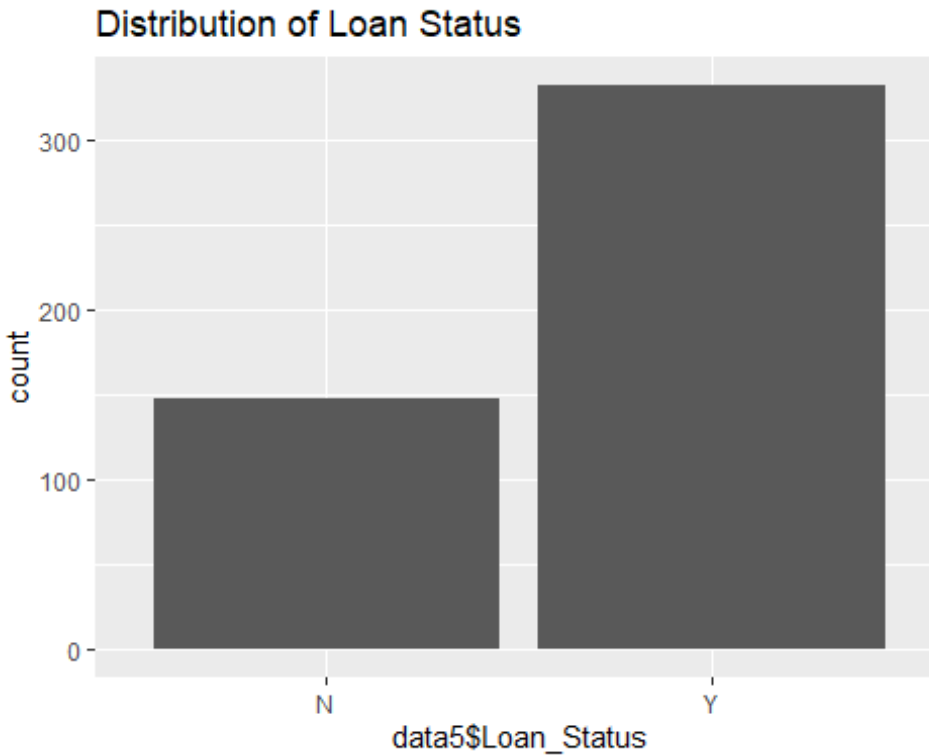
```
##
##      N   Y
## 0    87 187
## 1    28  52
## 2    20  65
## 3+    13  28
```

```
# Distribution of Loan status
```

```
ggplot(data5, aes(x= data5$Loan_Status))+geom_bar()+ggtitle("Distribution of
Loan Status")
```

```
## Warning: Use of `data5$Loan_Status` is discouraged.
```

```
## i Use `Loan_Status` instead.
```



STEP 4: DATA PREPROCESSING

#we first drop the loan id column income, ApplicantIncome and CoapplicantIncome since we no longer need them

```
data6 = data5[, !(names(data5) %in%
c("Loan_ID", "income", "ApplicantIncome", "CoapplicantIncome"))]
head(data6)
```

```
##   Gender Married Dependents   Education Self_Employed LoanAmount
## 1  Male    Yes           1   Graduate         No         128
## 2  Male    Yes           0   Graduate         Yes          66
## 3  Male    Yes           0 Not Graduate         No         120
## 4  Male    No            0   Graduate         No         141
## 5  Male    Yes           2   Graduate         Yes        267
## 6  Male    Yes           0 Not Graduate         No          95
```

```
##   Loan_Amount_Term Credit_History Property_Area Loan_Status
income_transformed
```

```
## 1           360           1      Rural      N
8.714568
## 2           360           1      Urban      Y
8.006368
## 3           360           1      Urban      Y
8.505323
## 4           360           1      Urban      Y
8.699515
## 5           360           1      Urban      Y
9.170872
```

```
## 6          360          1      Urban      Y
8.255569
```

```
# checking the number of observations
```

```
dim(data6)
```

```
## [1] 480  11
```

```
# dividing data into training (80%) and testing (20%)
```

```
train_split = data6[1:384,]
```

```
dim(train_split)
```

```
## [1] 384  11
```

```
test_split = data6[385:480,]
```

```
dim(test_split)
```

```
## [1] 96 11
```

```
# Apply feature scaling for numeric variables
```

```
scale_features = function(x){
  num_col = sapply(x, is.numeric)
  x[num_col] = scale(x[num_col])
  return(x)
}
```

```
scaled_train = scale_features(train_split)
```

```
scaled_test = scale_features(test_split)
```

STEP 5: MODEL FITTING We will compare different models and select the best one from 1. logistic regression 2. naive bayes 4. random forest

```
# first, we will use logistic regression
```

```
attach(scaled_train)
```

```
model = glm(Loan_Status ~ ., family = binomial(link = 'logit'), data =
scaled_train)
```

```
detach(scaled_train)
```

```
# printing the summary statistics for our logistic model
```

```
summary(model)
```

```
##
```

```
## Call:
```

```
## glm(formula = Loan_Status ~ ., family = binomial(link = "logit"),
##      data = scaled_train)
```

```
##
```

```
## Deviance Residuals:
```

```
##      Min       1Q   Median       3Q      Max
## -2.2654  -0.4230   0.4547   0.6942   2.3520
```

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error z value Pr(>|z|)
```

```

## (Intercept)          -3.05921    0.59867   -5.110 3.22e-07 ***
## GenderMale           0.44288    0.37233    1.189 0.23425
## MarriedYes           0.37696    0.32423    1.163 0.24498
## Dependents1          -0.66458    0.37989   -1.749 0.08022 .
## Dependents2           0.37682    0.43653    0.863 0.38802
## Dependents3+         0.26450    0.54571    0.485 0.62790
## EducationNot Graduate -0.38808    0.35871   -1.082 0.27931
## Self_EmployedYes     -0.44277    0.41806   -1.059 0.28955
## LoanAmount           -0.57555    0.20238   -2.844 0.00446 **
## Loan_Amount_Term     -0.28462    0.17170   -1.658 0.09738 .
## Credit_History1       3.62633    0.48590    7.463 8.45e-14 ***
## Property_AreaSemiurban 1.08703    0.34975    3.108 0.00188 **
## Property_AreaUrban    -0.00508    0.33902   -0.015 0.98805
## income_transformed    0.33122    0.21862    1.515 0.12975
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##    Null deviance: 478.56  on 383  degrees of freedom
## Residual deviance: 342.87  on 370  degrees of freedom
## AIC: 370.87
##
## Number of Fisher Scoring iterations: 5

```

Lets explain the results; 1. We can see that, only the Intercept, LoanAmount, Credit_History1 and Property_AreaSemiurban are significant in predicting the probability of a loan to be approved due to their loww p-value (<0.05). The main significant variable is Credit_History1 showing that having a credit history strongly increases the likelihood of loan approval.

2.The intercept (-5.710799) is the baseline log-odds of loan approval when all the predictors are set to 0.

3. At the top, we have a summary of the devince residual which is a measure of goodness of fit for the model. Since they are close to being centered on, they look symmetrical
4. The coeeficient for the GenderMale (0,442876) shows that, being male increases the log-odds of loan approval compare to the baseline category (female)
5. For education, not being a graduate reduces the log-odd of loan approval by 0.388078.

Based on the above explanations, the other variables applies the same explanations compared to the respective baseline categories.

The null deviance is the deviance of the model with only the intercept The residual deviance is the model with the predictors included The AIC (Akaike Information Criterion) is the is a measure of the model quality.

```

# performing a chi-square tets for the model
anova(model, test = "Chisq")

## Analysis of Deviance Table
##
## Model: binomial, link: logit
##
## Response: Loan_Status
##
## Terms added sequentially (first to last)
##
##
##              Df Deviance Resid. Df Resid. Dev  Pr(>Chi)
## NULL                      383      478.56
## Gender                1      2.810      382      475.75 0.0936674 .
## Married                1      1.644      381      474.10 0.1997343
## Dependents            3      6.221      378      467.88 0.1013352
## Education              1      2.266      377      465.62 0.1322418
## Self_Employed          1      0.773      376      464.84 0.3792810
## LoanAmount             1      9.398      375      455.45 0.0021720 **
## Loan_Amount_Term       1      1.135      374      454.31 0.2867354
## Credit_History          1     94.589      373      359.72 < 2.2e-16 ***
## Property_Area          2     14.444      371      345.28 0.0007304 ***
## income_transformed     1      2.414      370      342.87 0.1202912
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

# now lets apply the fitted model to the test dataset
fitted.results = predict(model, scaled_train, type = 'response')
fitted.results = ifelse(fitted.results > 0.5,1,0)
misClasificError = mean(fitted.results != scaled_train$Loan_Status)
head(fitted.results)

## 1 2 3 4 5 6
## 1 1 1 1 1 1

```

STEP 5: This the last step it involves plotting the ROC curve and calculating the AUC.

TO BE CONTINUED.....